



AI Internship Challenge – Text Classification Task

Goal

Build a small AI/ML project to demonstrate your ability to work with data, preprocess it, train a simple model, and evaluate results.

Dataset

Choose any **publicly available text classification dataset** (e.g., news, spam detection, sentiment analysis).

Requirements:

- At least **2–5 categories** of text.
- Enough samples to train and test a model (minimum a few thousand entries).
- The dataset must be publicly accessible.

Include the dataset link/source in your [README .md](#).

Tasks / Features

Data Preparation

- Load and inspect the dataset.
- Perform basic preprocessing (clean text, remove stopwords, lowercase, etc.).

Exploratory Analysis

- Show simple statistics (e.g., number of samples per category, most frequent words).

Model Training

- Train a simple classifier (e.g., Logistic Regression, Naive Bayes, or a small Neural Network)
- Split the dataset into training and test sets
- Evaluate model accuracy (optionally include precision, recall, F1-score)

Prediction Script

- Provide a function or script where a user can input a sentence and the model predicts its category

Bonus

- Use TF-IDF or embeddings instead of raw bag-of-words
- Compare at least two models and explain which performs better
- Create a small visualization (e.g., confusion matrix)
- Wrap the prediction script in a simple Flask or FastAPI endpoint

Submission Instructions

- Push your code to a public GitHub repository
- Include a short README.md with instructions and explanations
- After completing, send the repository link to **office@linkplus-it.com**
- Any commits with timestamps after sending the confirmation email will **not** be considered